



# Advanced Topics in Dimensional Modeling

**DAMA Philadelphia**  
**January 8, 2008**

**Presented by**  
**Tom Haughey, President**  
**InfoModel LLC**  
**868 Woodfield Road**  
**Franklin Lakes, NJ 07417**  
**USA**  
**201 755 3350 ©**  
**201 337 9094**  
**[tom.haughey@InfoModelUSA.com](mailto:tom.haughey@InfoModelUSA.com)**






# What Is A Dimensional Model?

---

- A dimensional model is a model in which the data is structurally classified as fact or dimension.
- General characteristics:
  - Query oriented
  - Structured around data usage not business rules
  - Organized roughly into base facts and dimensions of those facts
  - Based on identification of key grains of data and on characteristics of those grains
  - Consisting usually of snapshot, conglomerated data
  - Looks to reduce the number and depth of joins
- Two general patterns-
  - **Star** schema in which all multi-leveled dimensions are flattened
  - **Snowflake** in which at least one multi-leveled dimension is kept separate





# Star or Snowflake?

---

- First, some technologies require a snowflake and others require a star.
- Second, some queries naturally lend themselves to a breakdown into fact and dimension. Not all do. Where they do, a star is generally a better choice.
- Third, there are some business requirements that just cannot be represented in a star.
- Fourth, a snowflake should be used wherever you need greater flexibility in the interrelationship across dimension levels and components.
- Fifth, whether you collapse Header and Line Item into one fact table (and thereby reduce the structure to a star) should be based on tangible factors and by conformity to the dimensional pattern.
- Sixth, sometimes dimensional data changes at different rates or ways; in such more complex history situations, a snowflake can be better.
- Seventh, whether the star schema is more understandable than the snowflake is entirely subjective and anything but a foregone conclusion.
- Eighth, some tests have revealed no difference between the performance of a star and snowflake, and sometimes the snowflake was slightly faster.
- It is unwise to pre-determine what is the best solution. A number of important factors come into play and these need to be considered.






# Snowflake Caveat

---

- Remember
  - Even if you choose to snowflake, it is not necessary to snowflake EVERYTHING
  - A hybrid is often practical
  - Because a Snowflake requires a join (say) up a hierarchy, it does not necessarily mean it is an I/O
  - DBMSs are smart and do three things:
    - Fetch data in pages, not rows
    - Read ahead
    - Buffer data that is frequently used
  - The most expensive DBMS operations are not joins but sorts and Cartesian products
    - The volume of data in the join is more significant than the presence of a join






# Observations on the Dimensional Model

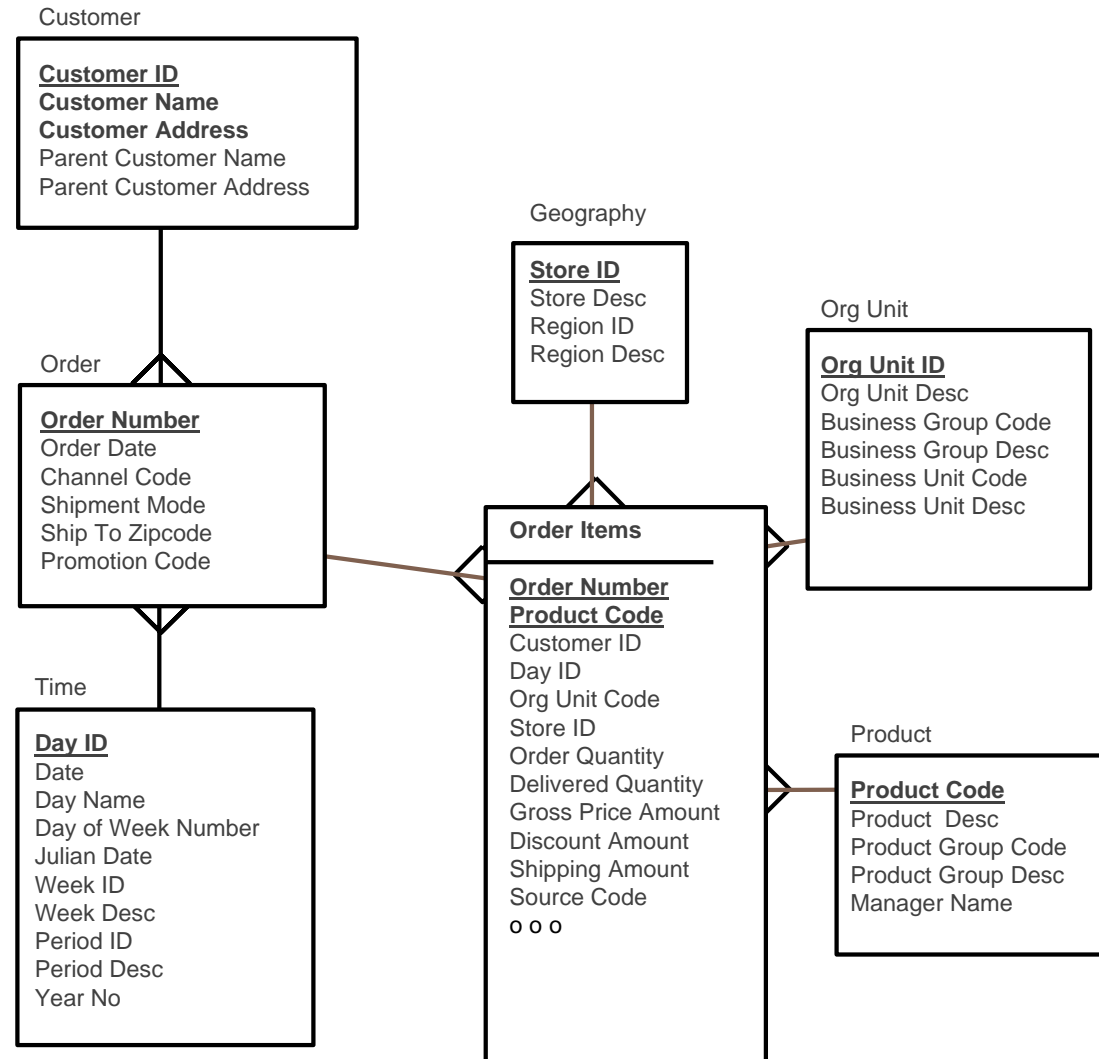
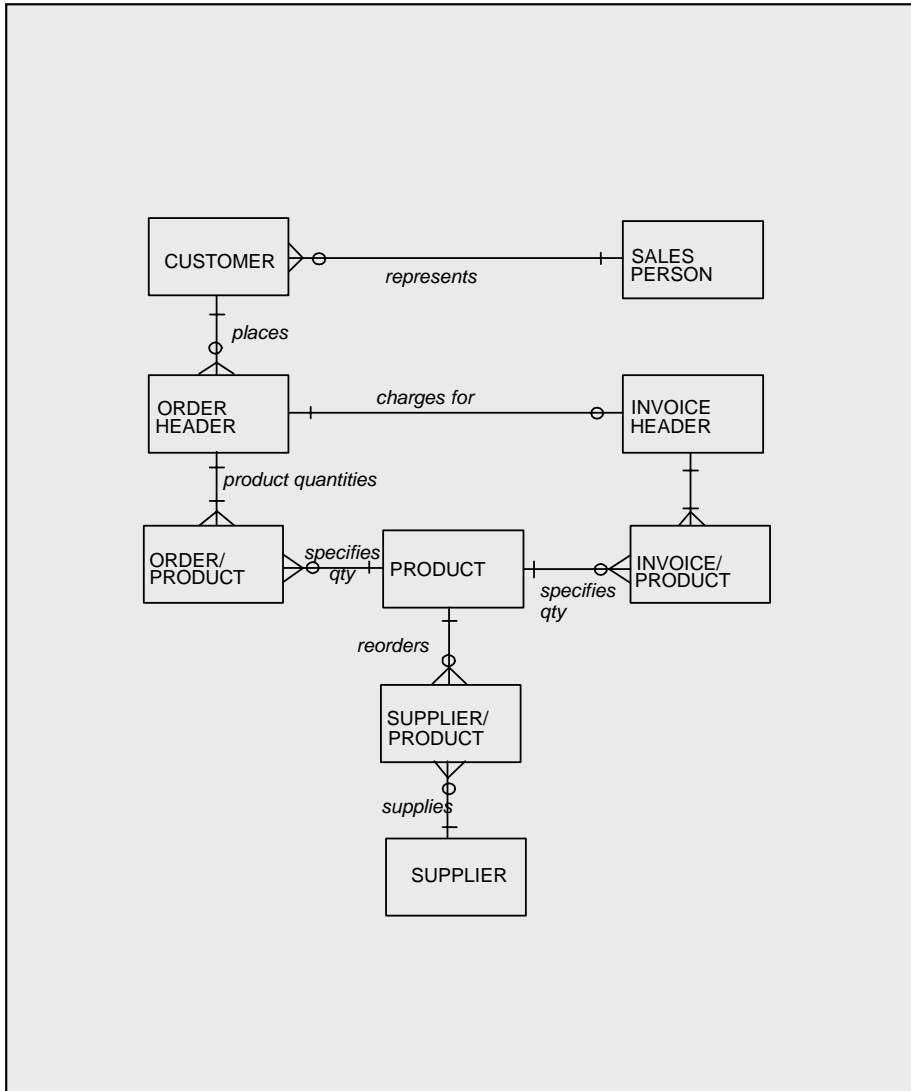
---

- The following points represent the perspective of this presentation and are deviations from classical dimensional modeling theory:
  - Inherently, data is neither a fact or dimension. This is determined by query usage and should not be explicitly declared
  - Queries do not have to access from dimension to fact, or v.v.
    - It can be meaningful to query facts directly
    - It can be meaningful to query dimensions directly
  - Dimensions can be related to and join other dimensions without going through a fact
  - Facts can be joined to other facts
  - A foreign key in a fact can be null (in any valid form of null representation) if the relationship is optional
  - The choice of a design pattern (e.g., star) should be based on heuristics and not decided *a priori*
  - However, the capability to achieve these things can be limited by a given technology (e.g., MDDBMS\* or Redbrick)






# Operational vs. Analytical Model



Operational Model

Analytical Model



# Normalization and the Dimensional Model

---

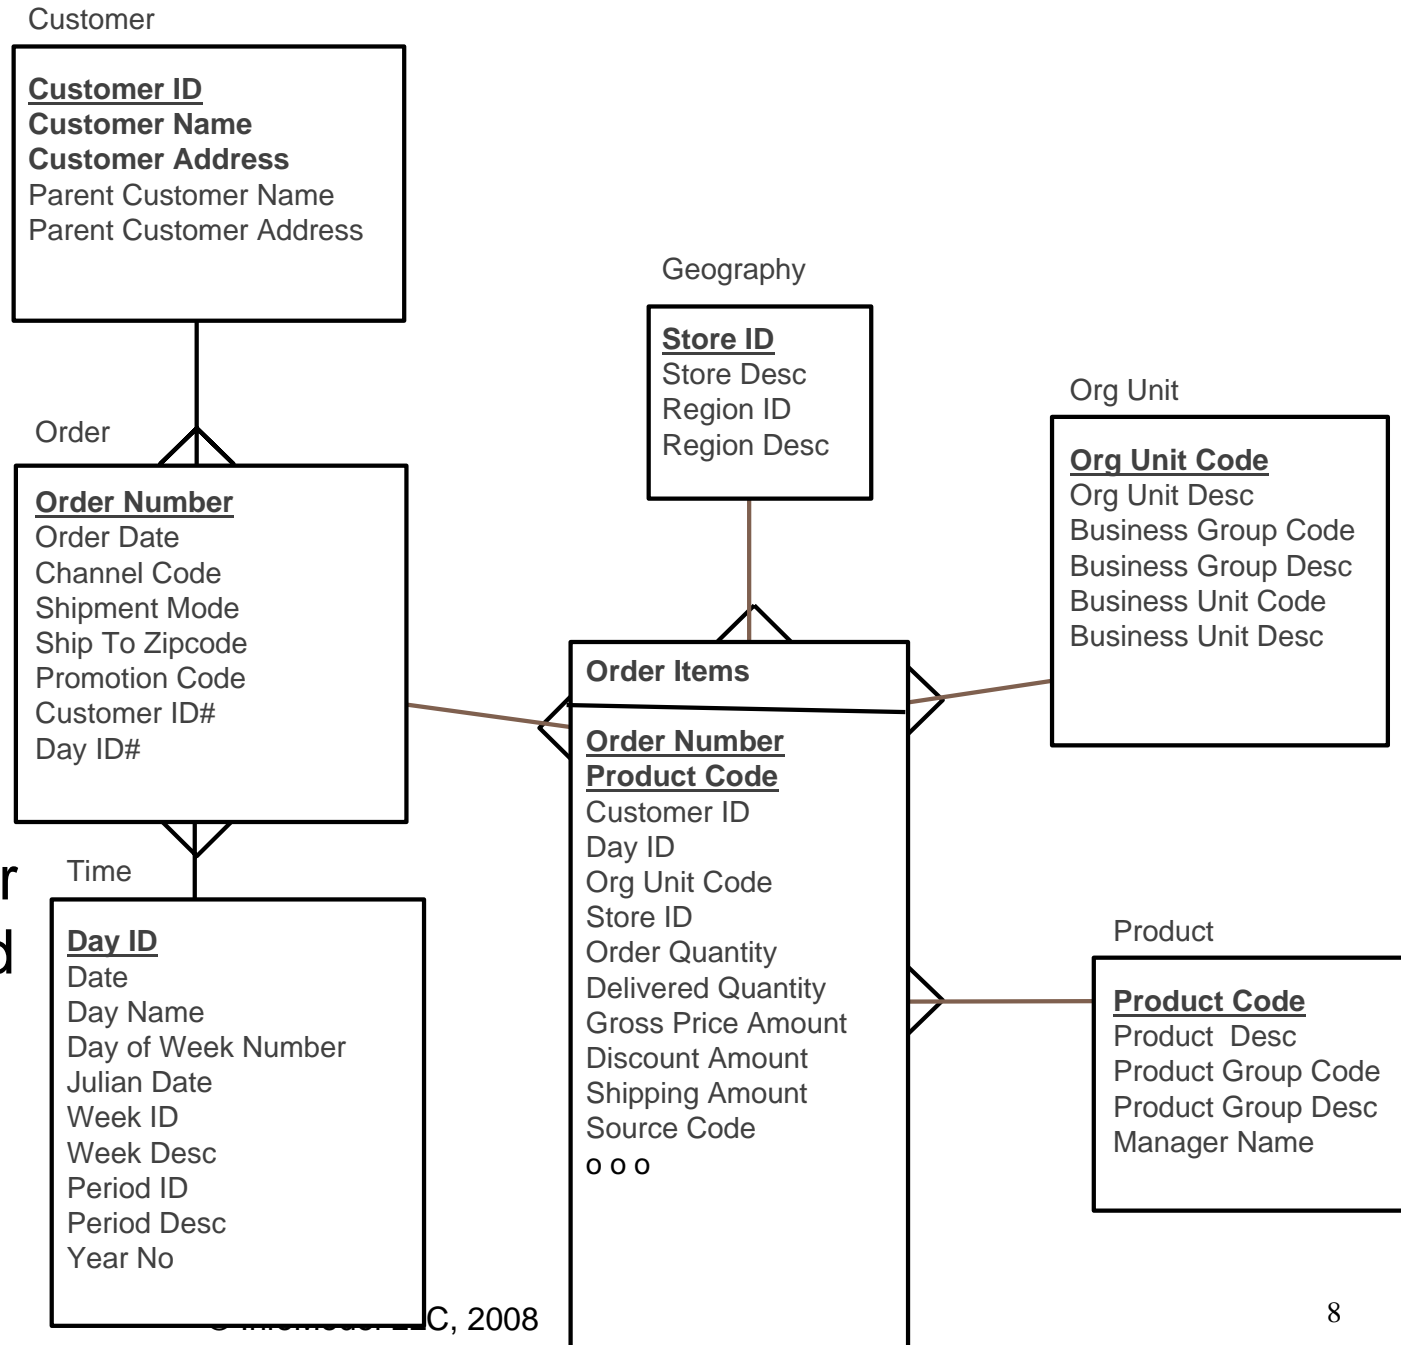
- Remember, if you take reporting or analytical data and normalize it, it will come out with a dimensional structure except for the following points:
  - Dimensions will be snowflaked
  - Compound facts (i.e., facts with multiple items with the transaction) will be divided into common data (the header) and the individual product sale (the line item)
  - For balance forward businesses, such as banking and insurance, facts will be divided into base facts (such as account or policy) and periodic status






# Complex Transactions (Line Item)

- Should you collapse Order Header?
- Consider the following factors:
  - Number of columns
  - Ratio of header to line item and
  - Data usage

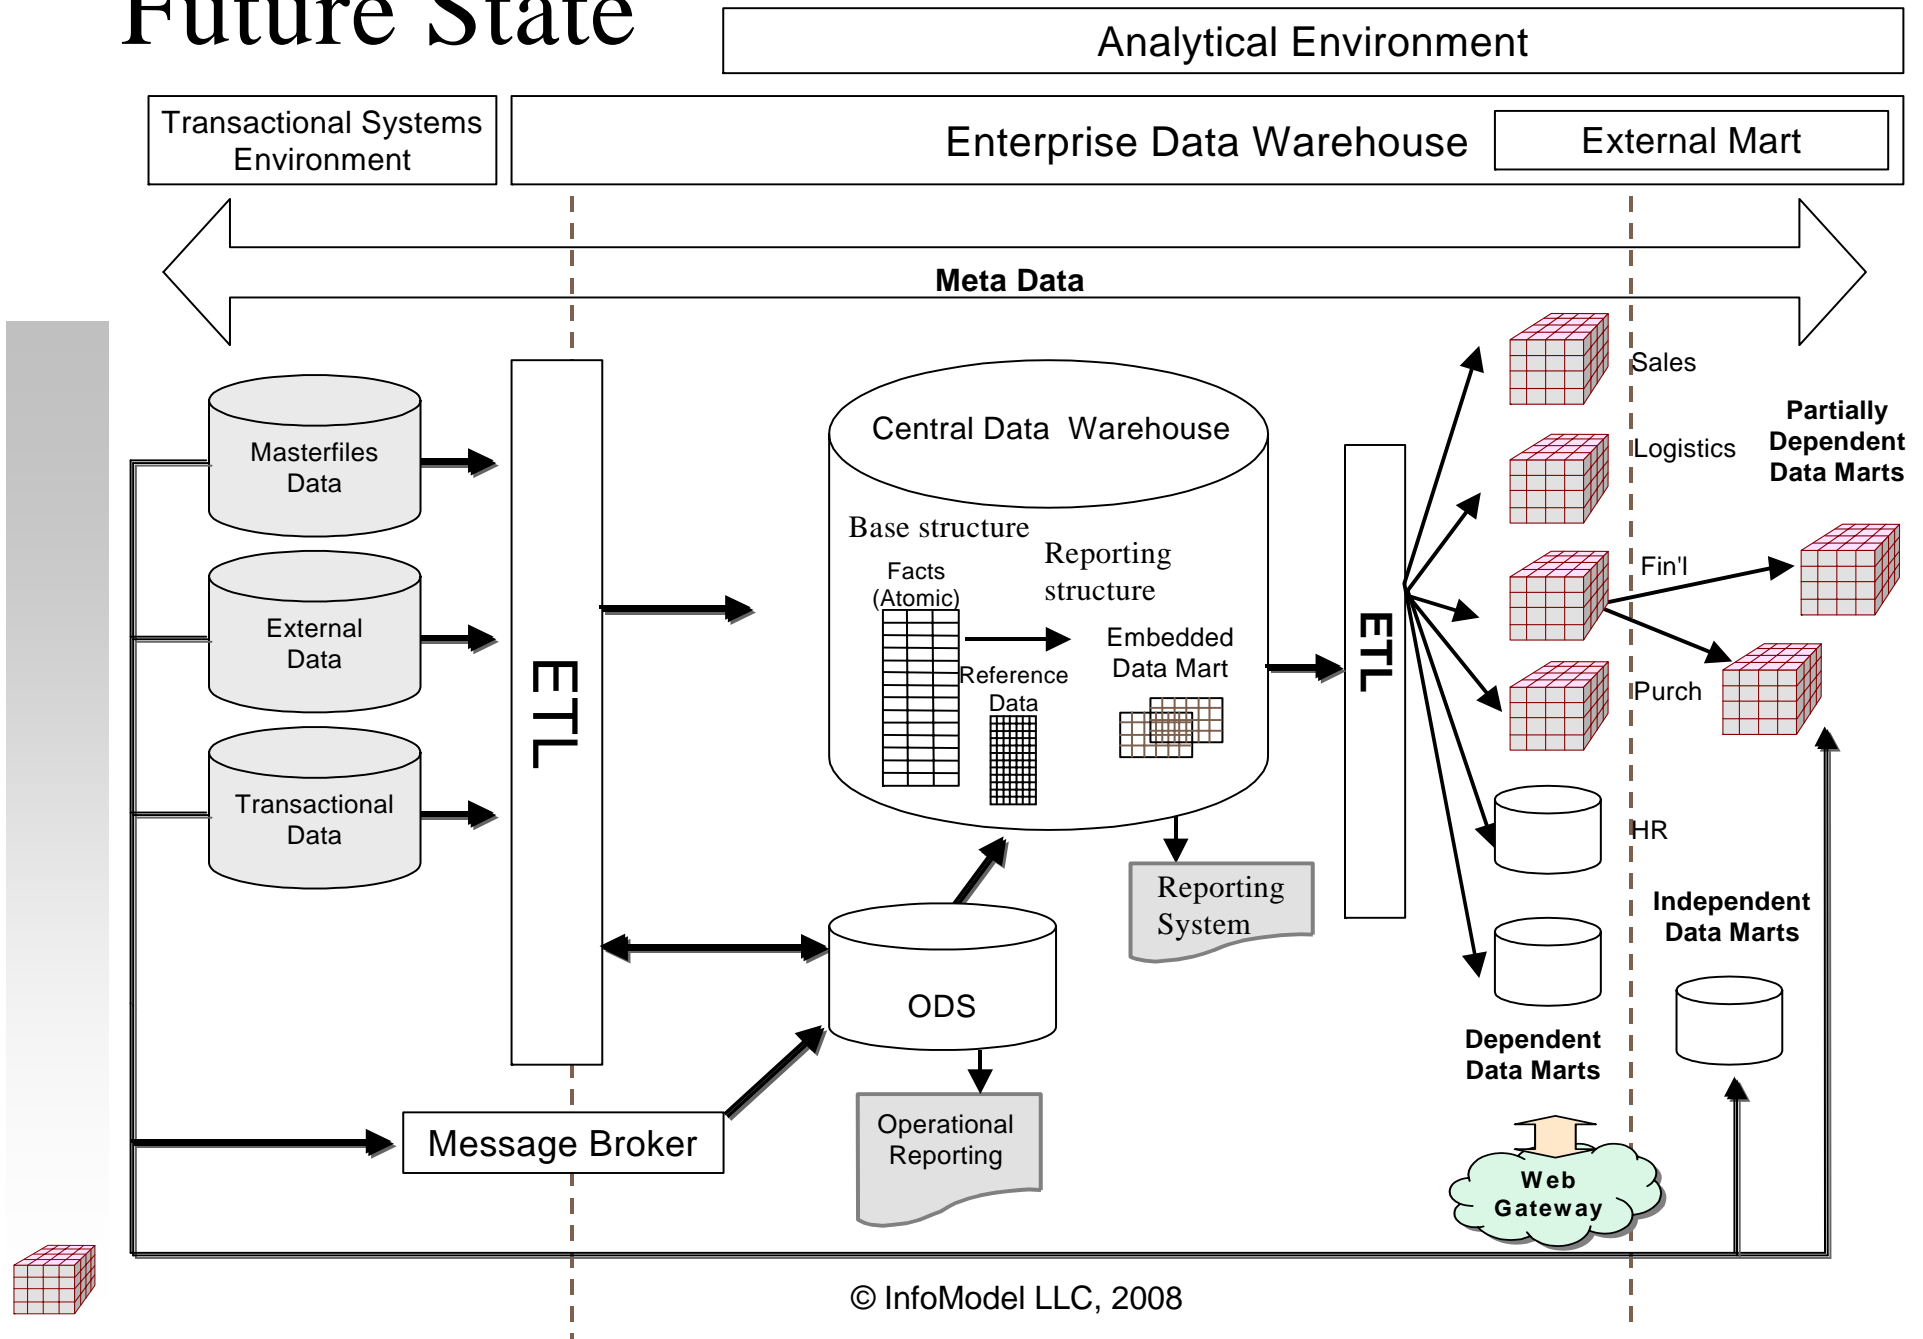






# The DW Environment

## Future State






# Overall Architecture of Data within the CDW

---

- **Base data:** DW data in its most atomic and flexible form; can be (and is) used for reporting
- **Reporting data:** data repackaged to facilitate reporting
  - **Detailed Reporting Data:** base data repackaged for reporting but of the same grain as base data, such as by flattening a recursive hierarchy
  - **Summary Reporting Data:** base data aggregated for reporting, derived from base data, or summarized in other ways





## Does the CDW Model Have to be Dimensional?

---

- The central data warehouse database (CDW) is the heart of the DW
- The CDW model has to satisfy information requirements and must do so within performance expectations.
  - It must use whatever data compromises are necessary to achieve this
  - Many factors play into this
    - All of them tangible
    - None of them emotional (such as having to fit a pattern)
  - Because of the vast amounts of data in the data warehouse, we will see that there are many levels of data and potentially many levels of optimization
  - For the CDW, it is best to start with a logical model of the information requirements
    - Business oriented
    - Independent of technology
    - Independent of implementation






# Functional Dependency

---

- Suppose technology were not an issue or that we had perfect technology
- The most flexible and open-ended model is one that honors functional dependencies (FD):
  - It shows the data with its most basic relationships
  - It shows the needs of the firm in its most basic form
  - FD's in operational data will be different than those in analytical data
    - What you are interested in will be different
    - The grains will be different
    - Business rules and rules-data will be different (such as pricing, discounting, etc.)
    - Either can be more or less detailed
    - Consider history in the DW
- The question is then:
  - Can your technology handle such a model?






## Is DM Best Suited for “Fixed” Queries?

---

- Dimensional models (DM) can be used to support different types of queries and the queries do not have to be predetermined
- However, DM does rely on the distinction between a fact and a dimension and in that sense its queries are fixed
- Only certain types of queries lend themselves to this pattern
- The following do not:
  - What 2 products are sold together most often and when?
  - What 2 products were sold together most often in Florida last year during hurricane season?
  - What are the main characteristics of customers who lapsed last month?
  - What is the most profitable combination of investments for this type of investment customer?






# Market Basket Analysis

---

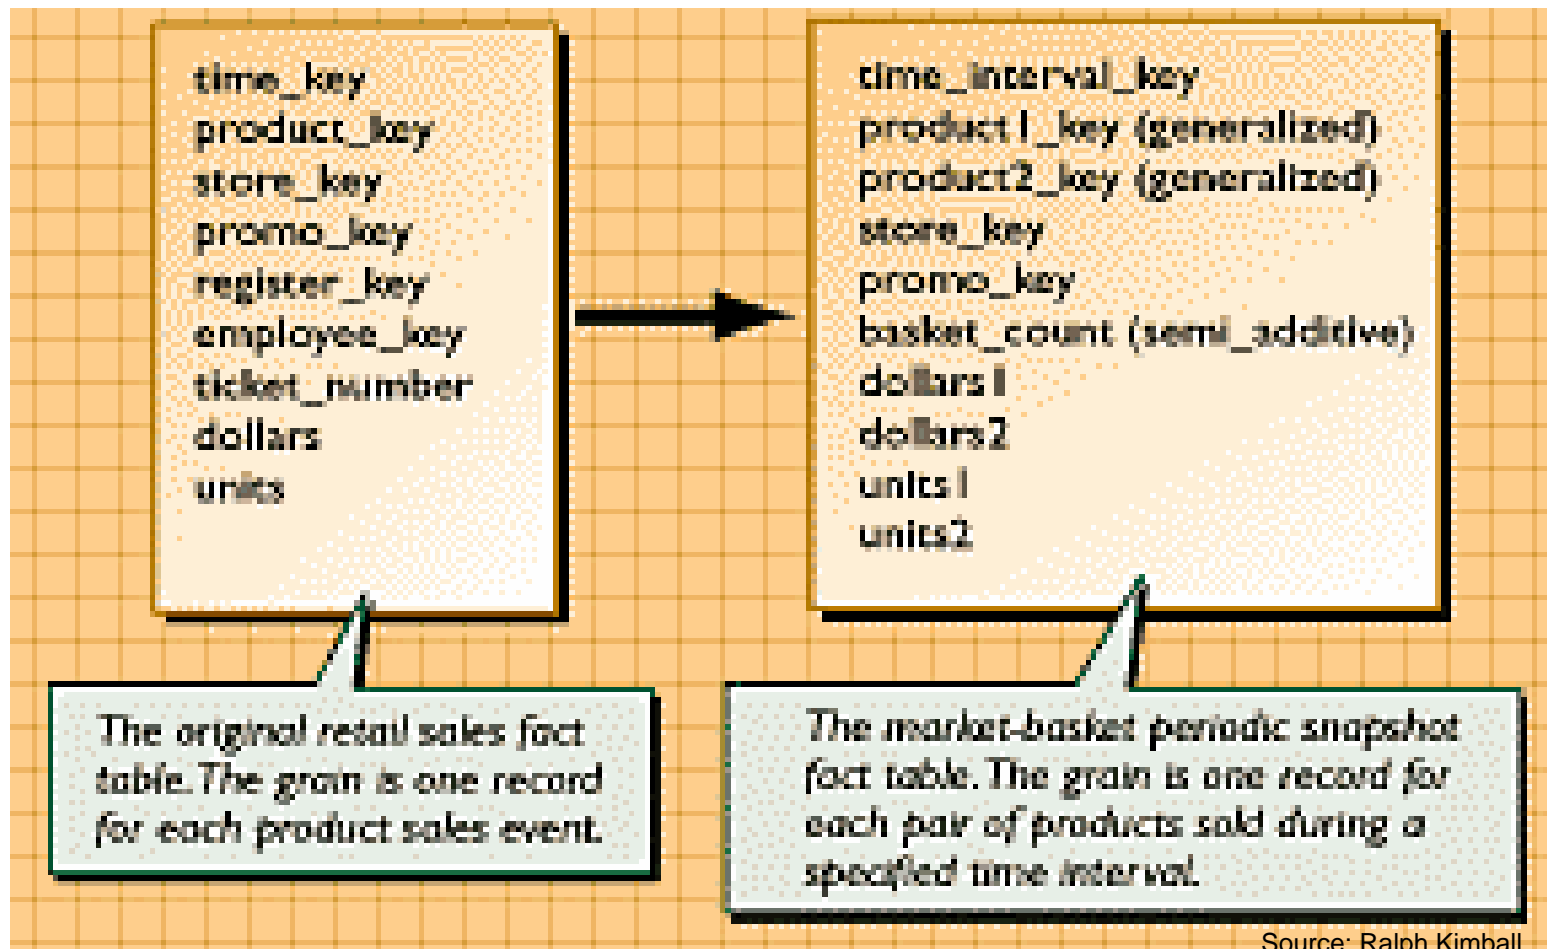
- Identifies what combinations of products sell together
- Can easily be applied to other situations:
  - What Basic Cable purchase are followed by Sports Cable purchases
  - What customer characteristics typically occur together in a given situation, such as policy cancellation
- Seeking to understand what meaningful product combinations are sold together in individual market baskets
- Applicable to:
  - Collocating products within store displays
  - Separating frequently combined products
  - Packaging and pricing
  - Understanding combinations of brands
  - Looking for mixed aggregate results more meaningful
  - Understanding what does and does not sell well together





## A Market Basket Example

- This example is well known but begs the question, meaning the statement presupposes the answer
- What do you do if you need the top **4** products, or in an insurance model, the top **6** characteristics of customers who lapsed last month? Or one time one thing, another time a another thing?






# Conformed Dimensions: Equality Rule

Customer

## Customer SK

Customer ID  
Customer Name  
Telephone No  
Contact Name  
....

**Demographics**



Marketing  
Data Mart 1


Conformed dimensions may have  
some individual attributes

Customer

## Customer SK

Customer ID  
Customer Name  
Telephone No  
Contact Name  
....

**Financials**



Financial  
Data Mart 2



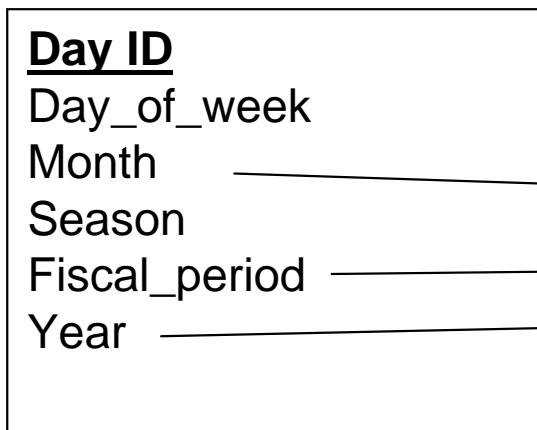





## Conformed Dimensions: Rollup Rule

- An attribute shared by multiple conformed dimensions must have the same business meaning and name so that it can be used as a common row header in separate queries
- Partially overlapping conformed dimensions are possible as the example shows:

### Calendar (Day)



### Calendar (Month)



- All data marts must deploy the conformed dimensions simultaneously
- All aggregates affected by a dimensions change must be removed or modified





# Conformed Dimensions

---


- Federation demands conformity
- Conformed dimensions means that when multiple copies of the same dimension exist, they are consistent with one another in that they honor either:
  - The equality rule
    - The same or compatible key
    - The same values for the key
    - The same grain or
  - The rollup rule
    - One is a strict rollup of the other
- Attributes in each conformed dimension can be a subset and can overlap
- Must be built at a common level of granularity (or be a rollup of the base dimension).
- Possibly, identified by a surrogate primary key.
- Conforming data in this way is productive, but it is difficult and costly
- Must establish an inviolable policy to reuse the conformed dimension whenever that subject is required for a data mart
  - Also called dependent dimensions or architected dimensions.





# Crisscrossed Facts

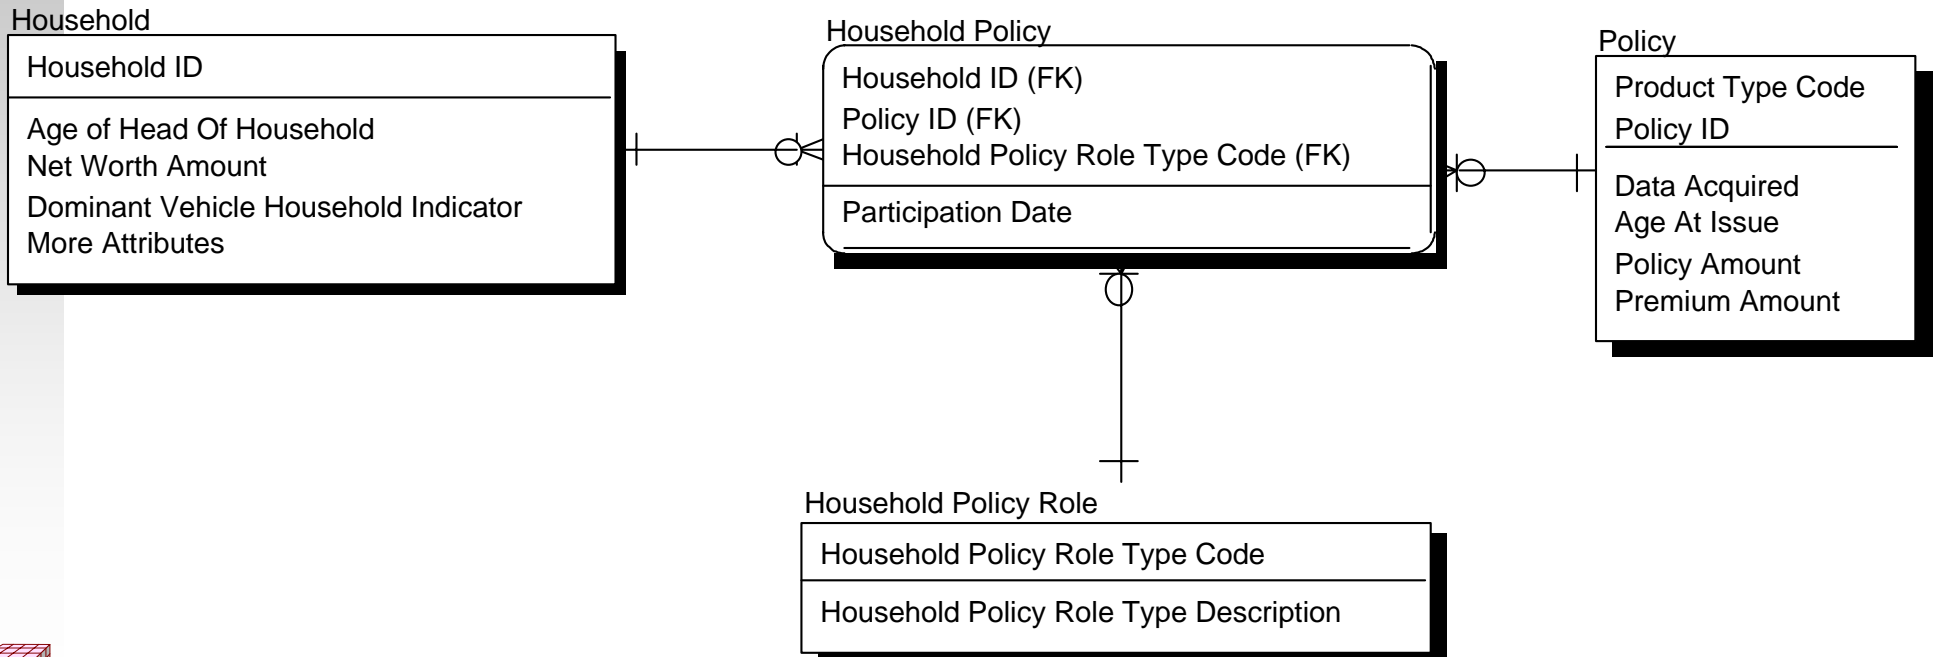
- Conformed dimensions but crisscrossed facts
- When DWs are decentralized, whether they are conformed or not, there is a natural tendency for one mart to need the data contained in another
- The problem occurs because facts are not shared across marts





# Snowflaked Many-to-Many Relationships

- Most dimension-to-fact relationships are one-to-many
- Some critical relationships can be many-to-many, such as Customer- or Household-to-Policy in insurance
- In the following example, assume Policy is the fact table






## Typical Multi-valued Dimensions

---

- Multi-valued dimensions have a M : M relationship to the fact
- Sometimes supported with a weighting factor or allocation factor
  - Individuals, households in an account or policy
  - Diagnoses for a patient treatment
  - Industry classifications (SICs) for a company
  - Segment classifications for a household
  - Other (and different) examples:
    - Large customer dimensions
    - Financial product dimensions
    - Multinational and decentralized enterprise calendar dimensions
  - Report additive measures both correctly allocated and overlapping

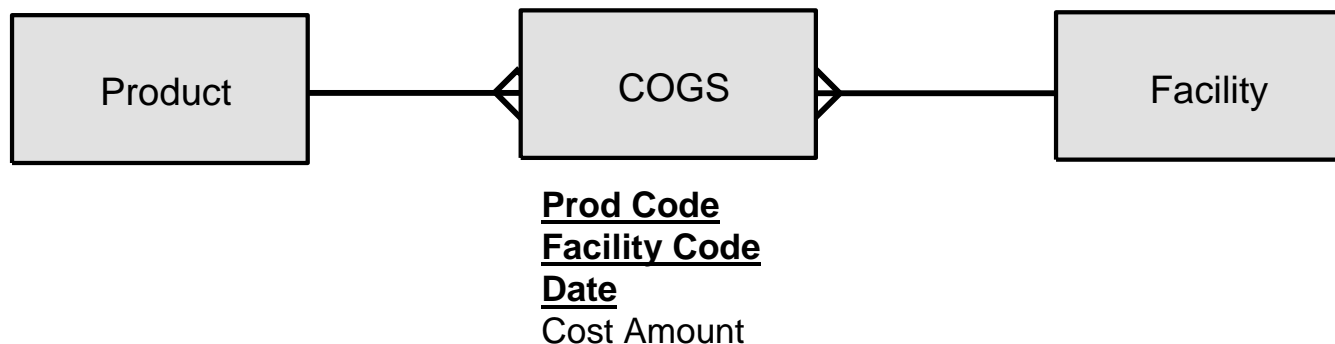





## Determining Facts or Dimensions

---

- Do all facts have to have a count or amount?
- Are all dimensions without them?
- Does it matter? Who cares?
- For example:
  - COGS (cost of goods sold – by product, time and location)
  - Factless facts
  - Give me the Customers that lapsed last month
  - Tell me which Facilities have Products with the lowest cost basis
  - Give me all Customers with Liquid Net Worth > \$2MM





## Periodic Status or Snapshot Requirements

---

- Often required in balance forward businesses such as banking and insurance
- Many transactions are not directly related to revenue
- Large number of transaction types, dimensionality and timing
- Revenue, status and other cumulative results needed at period end
- Achieved by the following:
  - Transaction time(s) rolled up to period end time
  - Transaction type dimension aggregated to periodic status dimension
  - Individual transaction amounts grouped and netted to multiple cumulative facts for period
  - Many other dimensions can remain the same





# Periodic Status or Snapshot

## Policy Transaction

### Transaction ID

Transaction DateTime ID  
Policy ID  
Product Type Code  
Customer ID  
Agent ID  
Coverage Code  
Covered Item Code  
Transaction Type Code  
Status Type Code  
Transaction Amount

## Periodic Status



### Period ID

### Policy ID

### Product Type Code

### Policy Status Code

Customer ID  
Agent ID  
Coverage Code  
Covered Item Code  
Status Type Code  
Premium Paid Amount  
Profitability Amount  
Face Value Amount  
Reserve Balance Amount  
Transaction Count

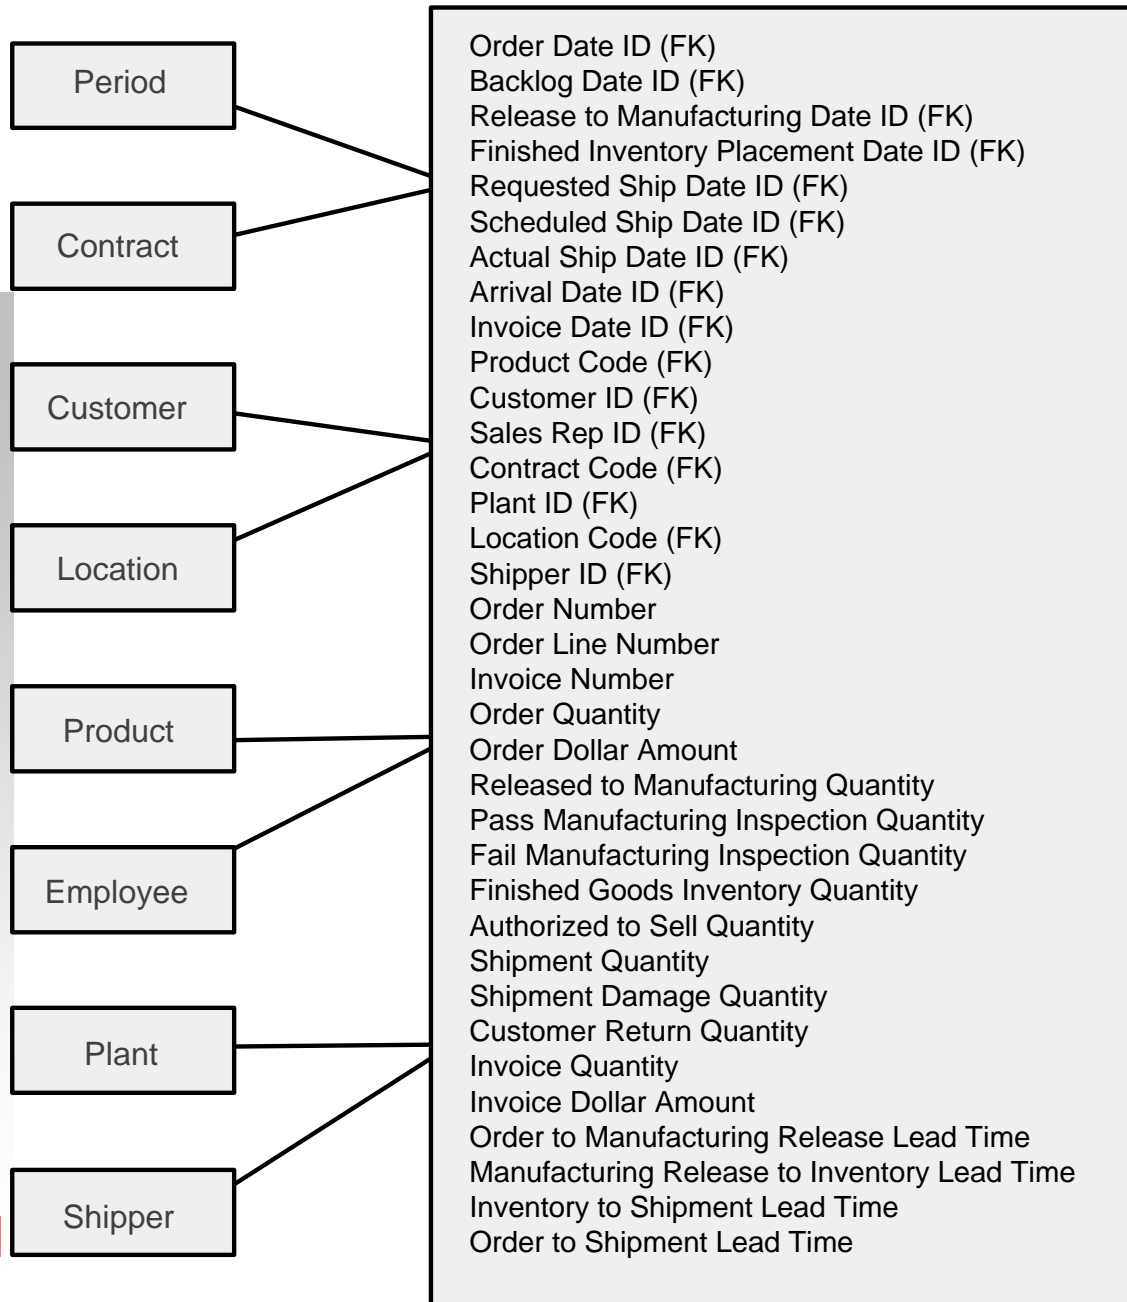






# Accumulating Snapshot

## Order Processing Accumulating Snapshot



- In this accumulating snapshot, the complete timeline of a process is represented.
- Each fact consists of multiple groups of data, including multiple dates.
- Each repeating group represents one stage of the process.

### • Pros

- All the data is collected in one place for easy retrieval

### • Cons

- Limit as to how much data you can reasonably carry for each repeating group
- Will work only if the relation of all stages is 1:1 to the order
- The row must be updated at each milestone of the process

### • Alternative


- One row for each milestone.






# Dimension-to-Dimension Joins


**1 : 1**



**1 : Many**




**Many-to-many**




- Dimension-to-fact relationships tell you what did happen
- Dimension-to-dimension relationships can tell you also what can happen
- Analyzing facts without dimensions is meaningless
- Analyzing dimensions, even without reference to facts, can be useful





# Fact-to-Fact Joins


- A relationship or join between two fact tables
- Considered taboo by dimensional modeling purists because:
  - It deviates from the dimension-fact paradigm
  - The table cannot be pre-classified as fact or dimension
  - Fact tables are voluminous requiring deep joins
  - Not supported by BI-specific technology especially:
    - Multidimensional DBMSs
    - Redbrick DBMS (only supports pure star schema)
    - However, data is data
- A robust DBMS platform will have no trouble accomplishing this join (often done via left outer join)





## Factless Facts

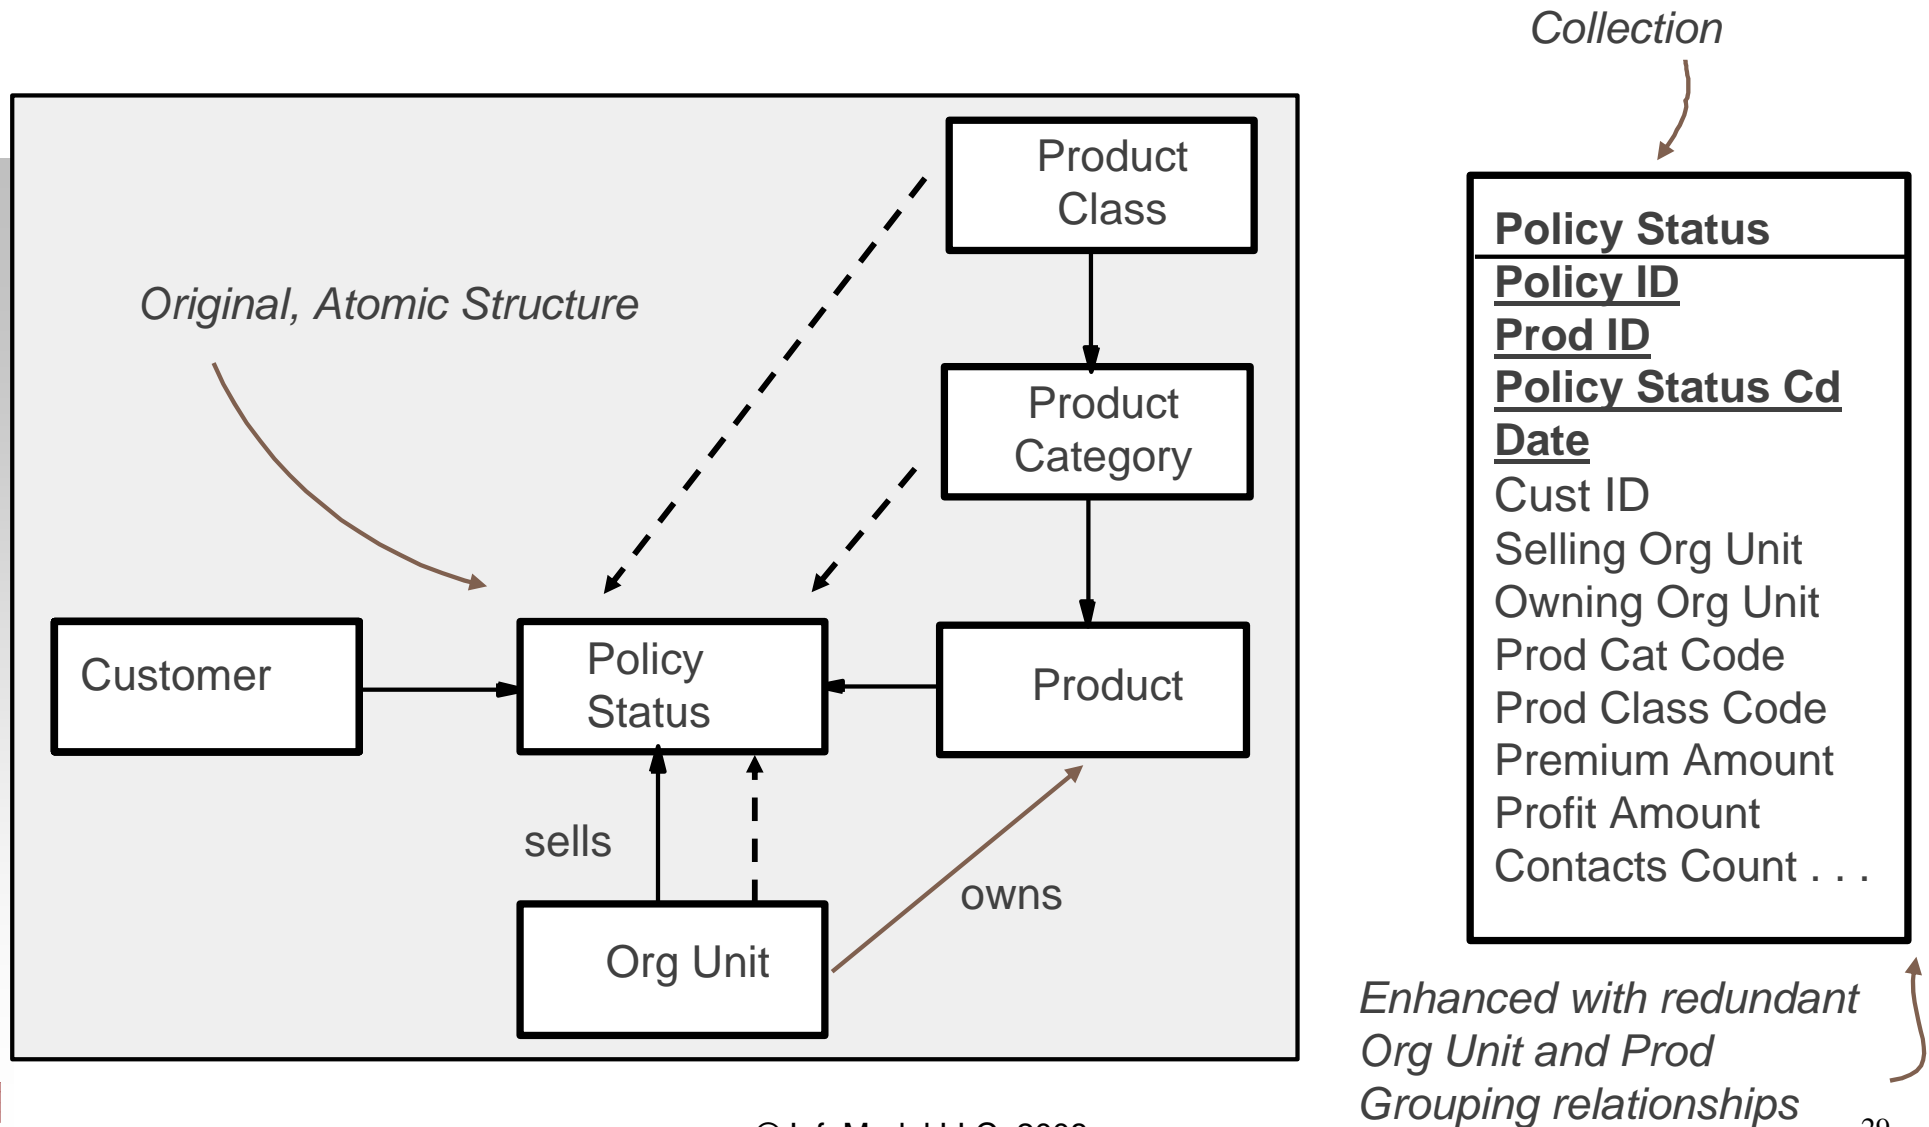
- Some “factless facts” may not be factless at all and could be expanded into collections of other data that can reasonably belong to it
- In effect, this is overloading the factless fact
  - Household - Household Policy - Policy
  - Customer - Customer Account - Account
    - Are Household Policy and Customer Account factless facts or complex dimensions? Before or after the changes suggested above?





# Collection


- Overloading a table by assembling into that table, the data or relationships that are often queried together, even though this may introduce some redundancy





## Very Large Dimensions (VLD)

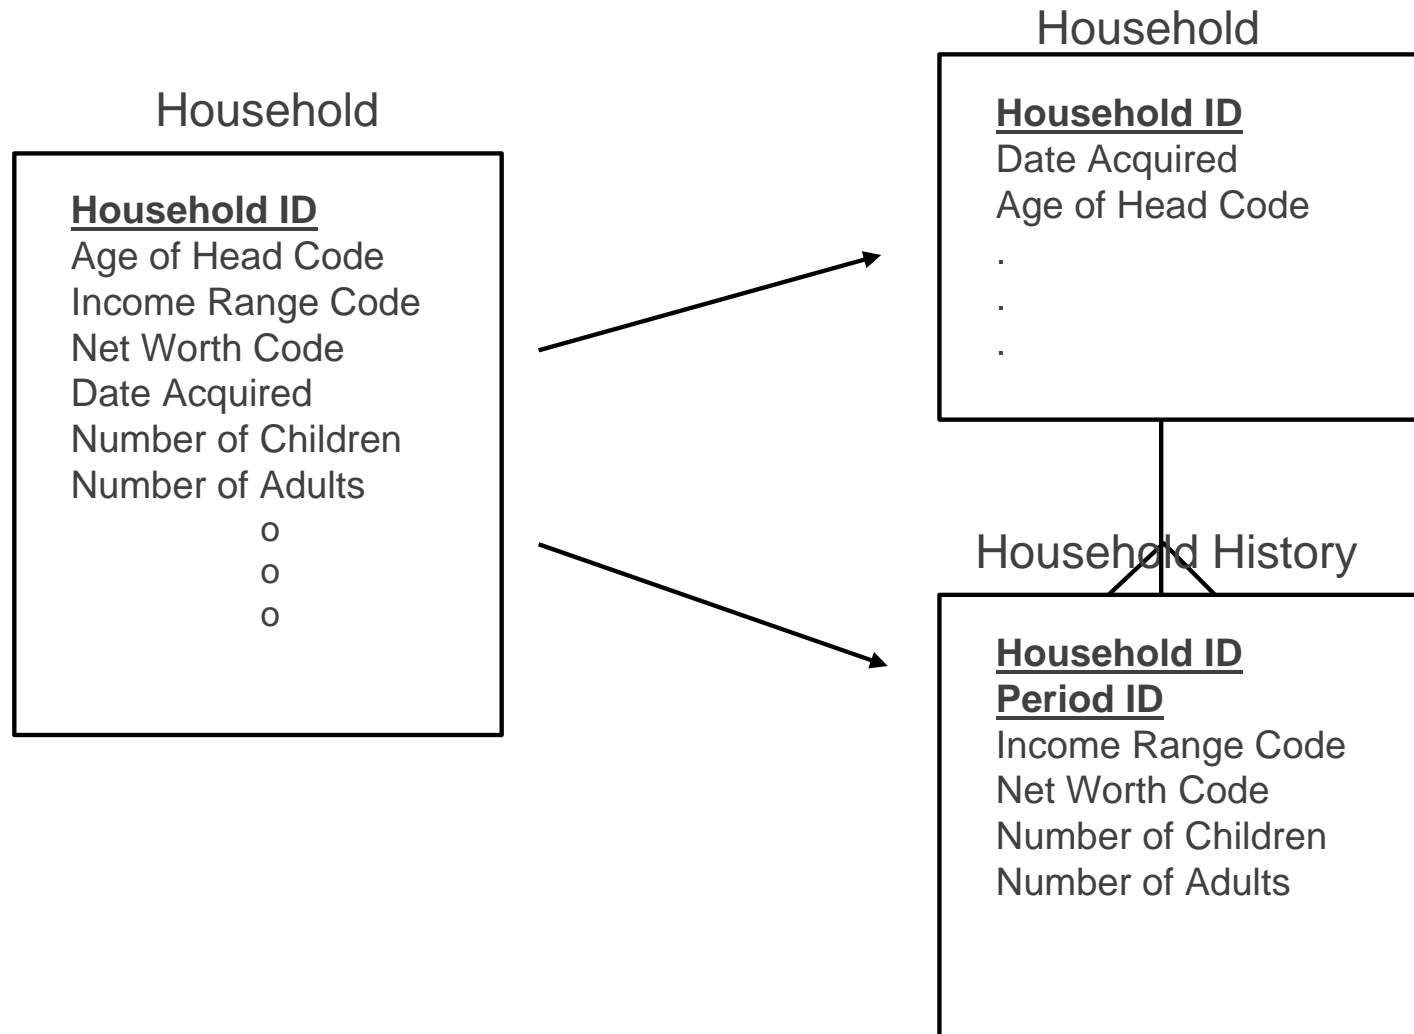
- Not all Dimension Tables are small compared with the size of the fact table
  - Examples of very large dimensions:
    - Customer Dimensions in banking, insurance, telephone companies, catalog retailers
    - Household dimensions in any retail business (e.g., insurance, brokerage)
- Can hold millions of records
- Supporting complex analytical processing requires both careful modeling as well as the use of sophisticated DSS indexing and join techniques.






# Very Large Dimensions (VLD)

- Split stable from changing data





# Heterogeneous Facts

---

- Heterogeneous facts are possible where the products offered or the markets (customers) addressed have different characteristics:
  - Heterogeneous customers:
    - Investment and insurance can have Private Clients and Capital Markets
  - Heterogeneous Products:
    - Again, insurance can have life, home, auto, etc., each with different attributes
    - Investment can have trades in stocks, options, etc., each with different attributes

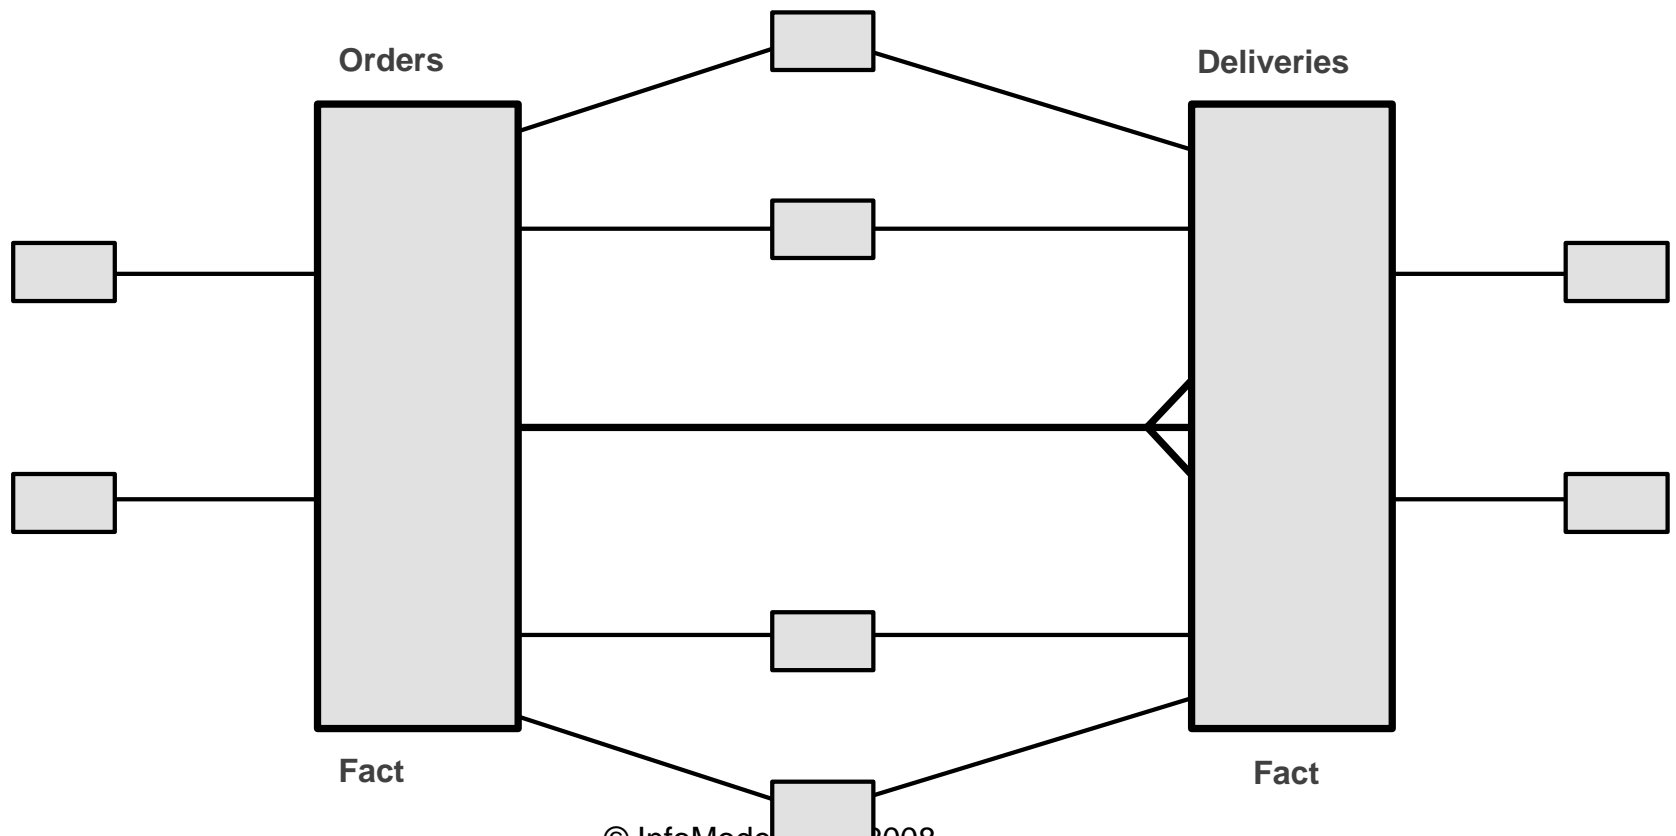






# Purely Heterogeneous Facts

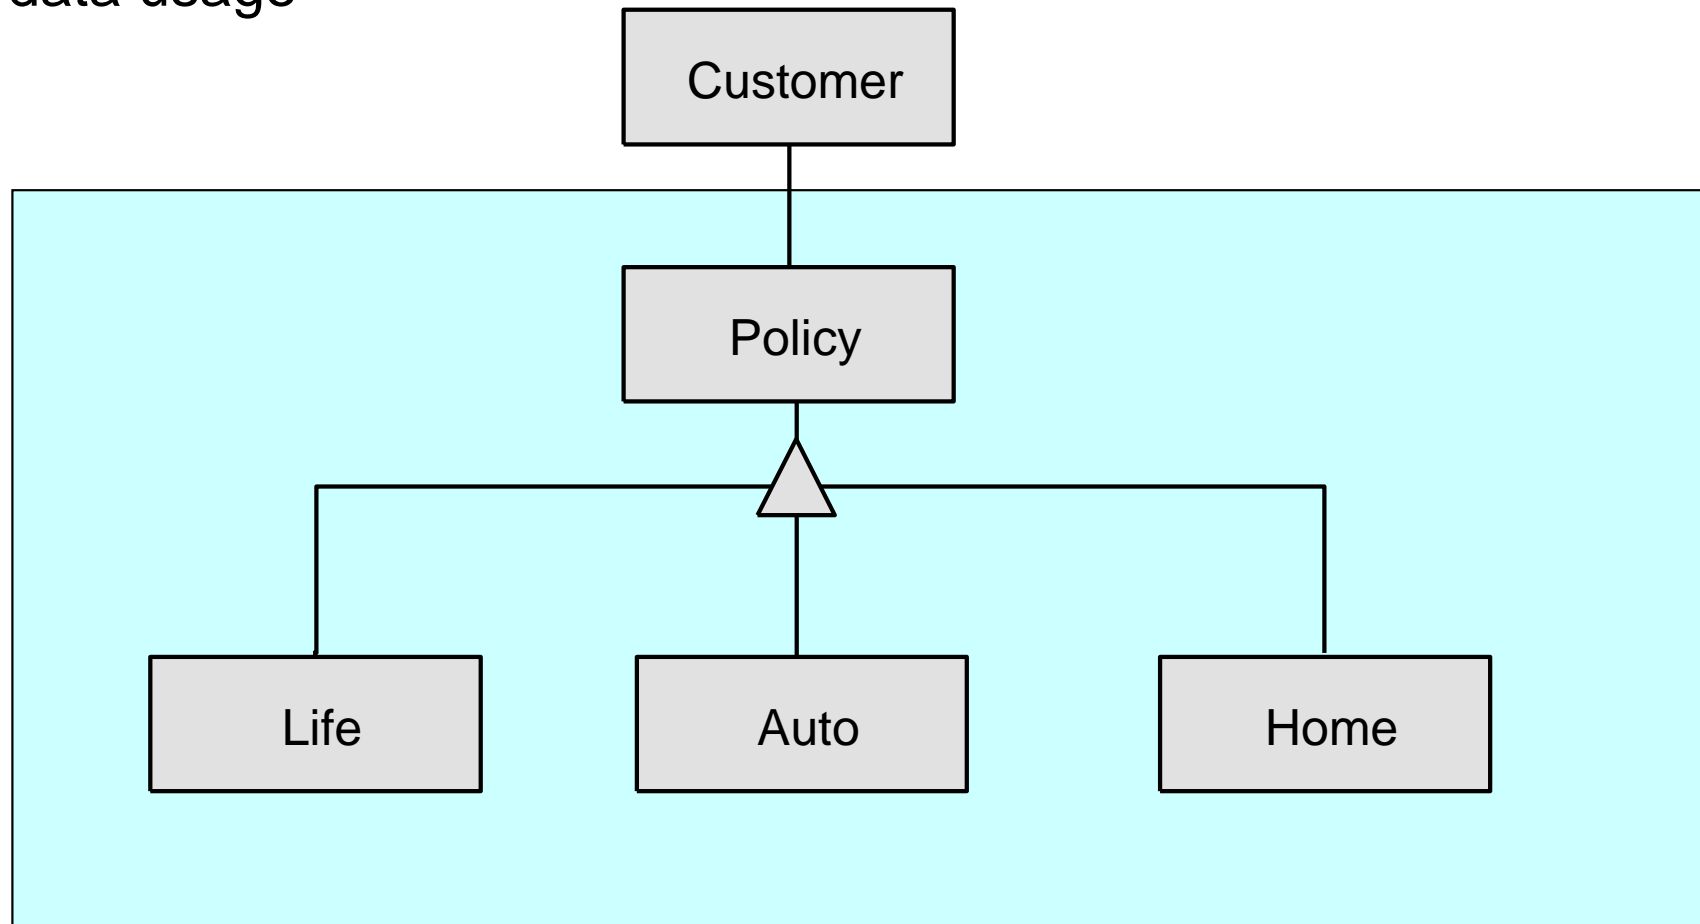
- If two facts are purely heterogeneous, such as, they have many different attributes or relationships, then they should simply be implemented as separate fact tables





## Subtype Heterogeneous Facts in LDM

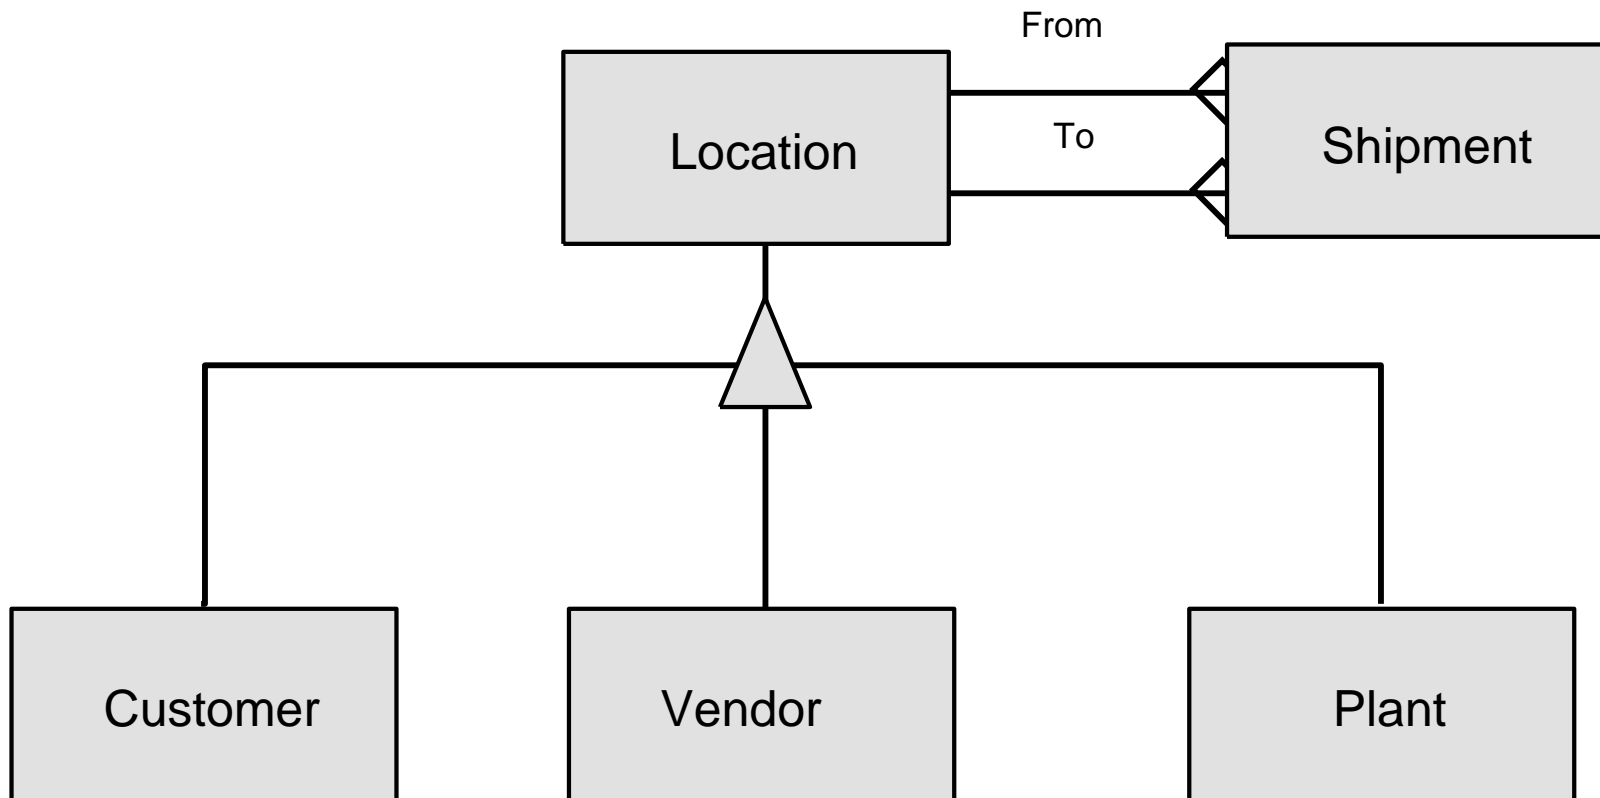
- When attributes or relationships are different, subtypes can be used
- Subtypes inherit all the characteristics of their supertype
- In implementing, consider number of columns, number of rows and data usage





## Supertype Common Roles

- The so-called “Unity” dimension pulls together different specializations that have common roles.






# Cartesian Dimensions

---

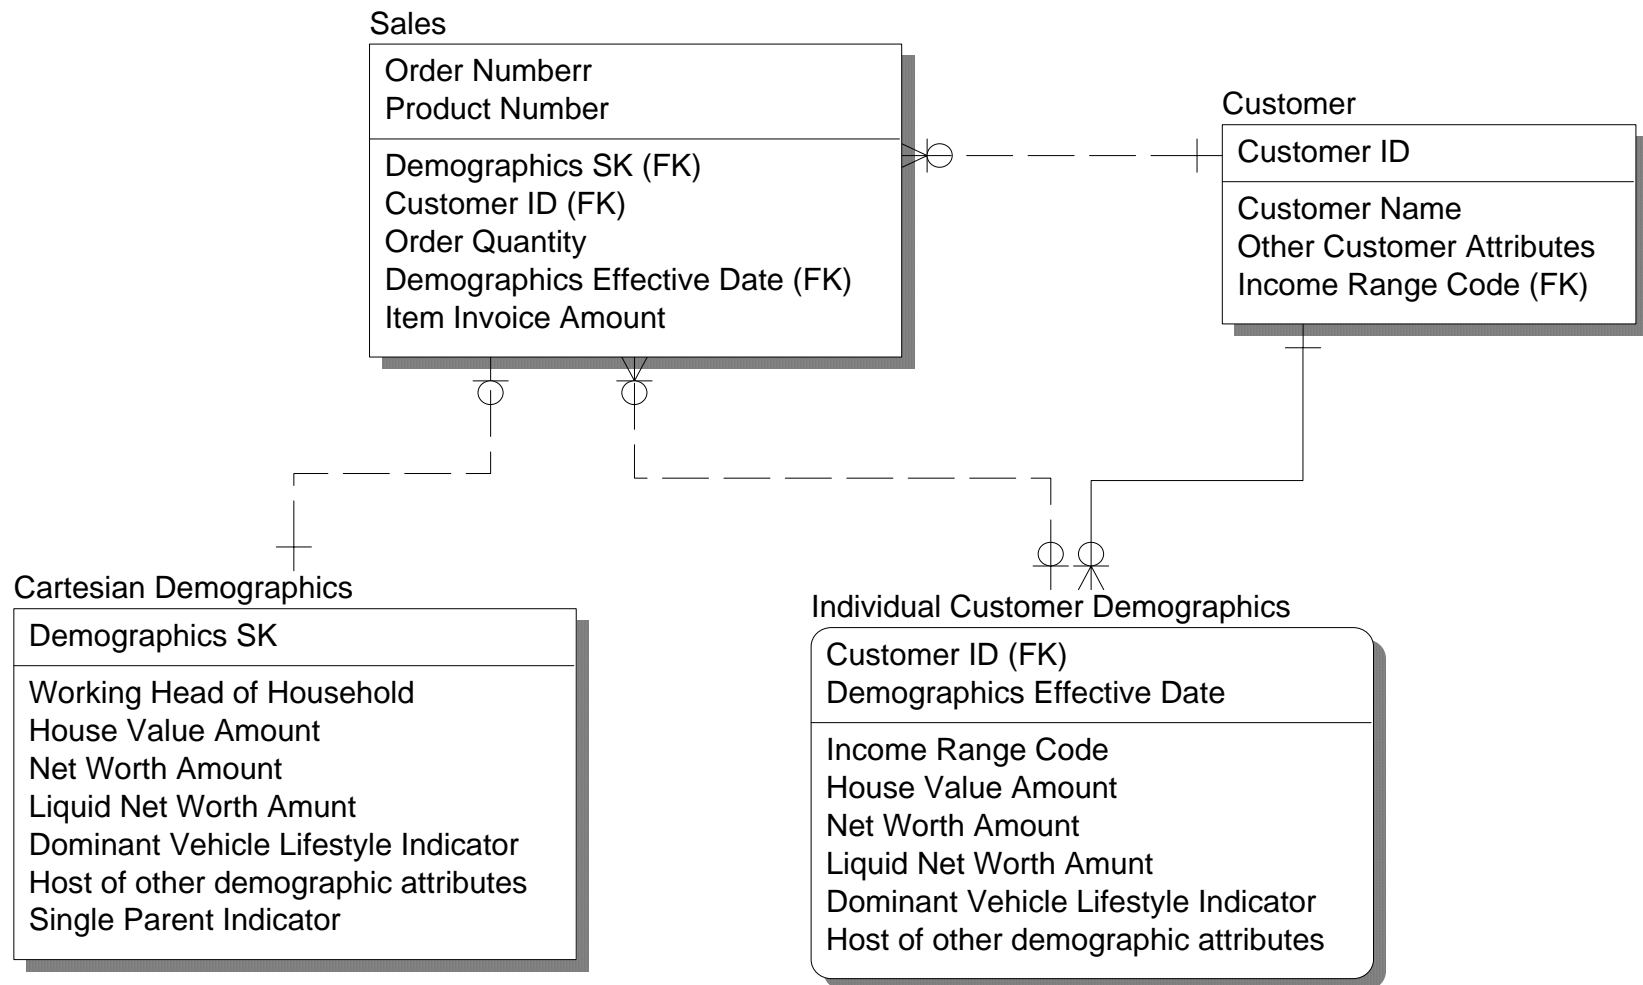
- Cartesian product is formed by combining all instances of all participating dimensions
- A composite dimension that contains all permutations of related characteristics such as demographics
- Created in advance
  - Related to parent entity
  - The alternative is to create instances in the composite dimension as they occur in the fact
  - Cartesian demographic makes retention of history more difficult to do





# Cartesian Dimensions

- For illustration purposes, this example shows both kinds of demographic dimensions related to the same fact
- Note surrogate key for Cartesian, history for Individual Customer Dimension







# Demographic Dimensions


- Can be constructed in two general ways:
  - Cartesian dimension: a dimension pre-populated with all possible combinations of demographic attributes
    - Can tell which combinations are not in use
  - Targeted dimensions: populated only when actual combinations exist
    - Cannot easily tell which combinations are not in use

**Cartesian  
Dimension**



**As-Used  
Dimension**






## Rapid Changes

---

- Data models should always FIRST represent the data by enforcing functional dependencies and THEN evaluate the need to deviate from that
- To store a dimensional attribute directly in the fact table should be done:
  - If the dimensions changes at the same rate as the fact occurs
  - If the history of the dimension is not preserved some other way such as via a dimension history

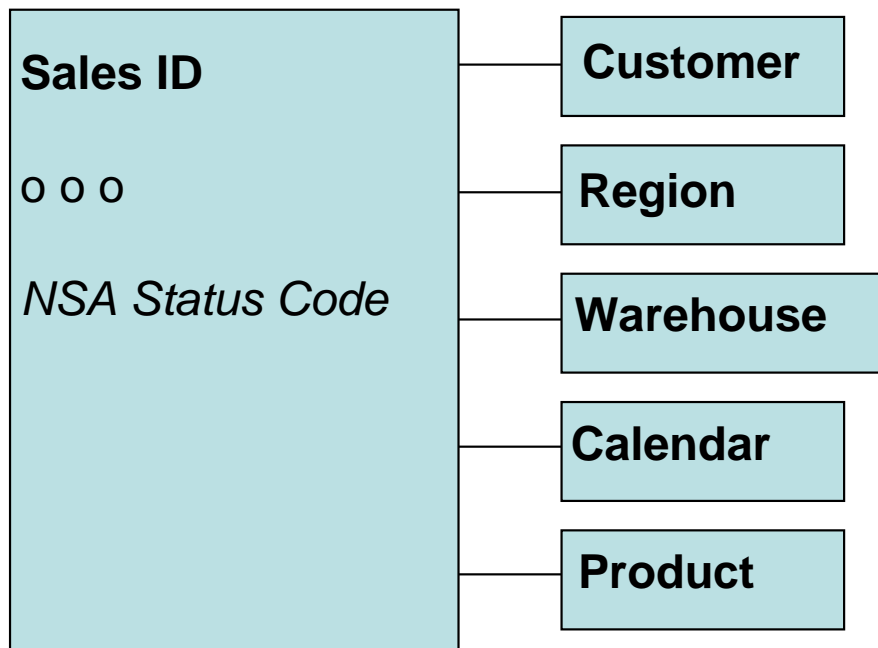





# Rapid Changes

- Conventional dimensional advice on rapidly changing dimensions (non-textual) is to put them in the fact table
- Which of the following is better for National Account Status (NAS)?

## Sales Fact



**Vs.**



- Be sure to examine overall data usage (of different types)
- Perhaps it is better to include NAS in both fact (to tell which sales were national account sales) and dimension (to tell which accounts are national accounts)







## Junk Dimensions

---

- A catch-all grouping of miscellaneous flags and indicators.
- Helpful, but not absolutely required, if a strong interrelationship among a group of miscellaneous dimensional attributes
- Used to reduce dimension clutter
- Collect together miscellaneous flags and indicators, and non-additive fact characteristics
- If possible, group such data that has a relationship
- Look throughout requirements for these associations
- Remove textual and unstructured data (such as comments) to its separate dimension





# Junk Dimension


---

- Benefits:
  - A useful location for related codes, indicators and their descriptors
  - Simplify a design that already has many dimensions.
  - Provide a smaller, quicker point of entry for queries
  - Capture the context of a specific transaction.
- Insurance example:
  - Capture the context surrounding claims.
  - Claims, even similar claims, may be handled differently.
  - E.g., how the claim was reported, investigated and paid
- Two approaches for creating junk dimensions:
  - Create in advance (if few and well known) . Each possible combination represents one row.
  - Create during ETL as instances are encountered (if many)





## Further Questions?





**Finis**

